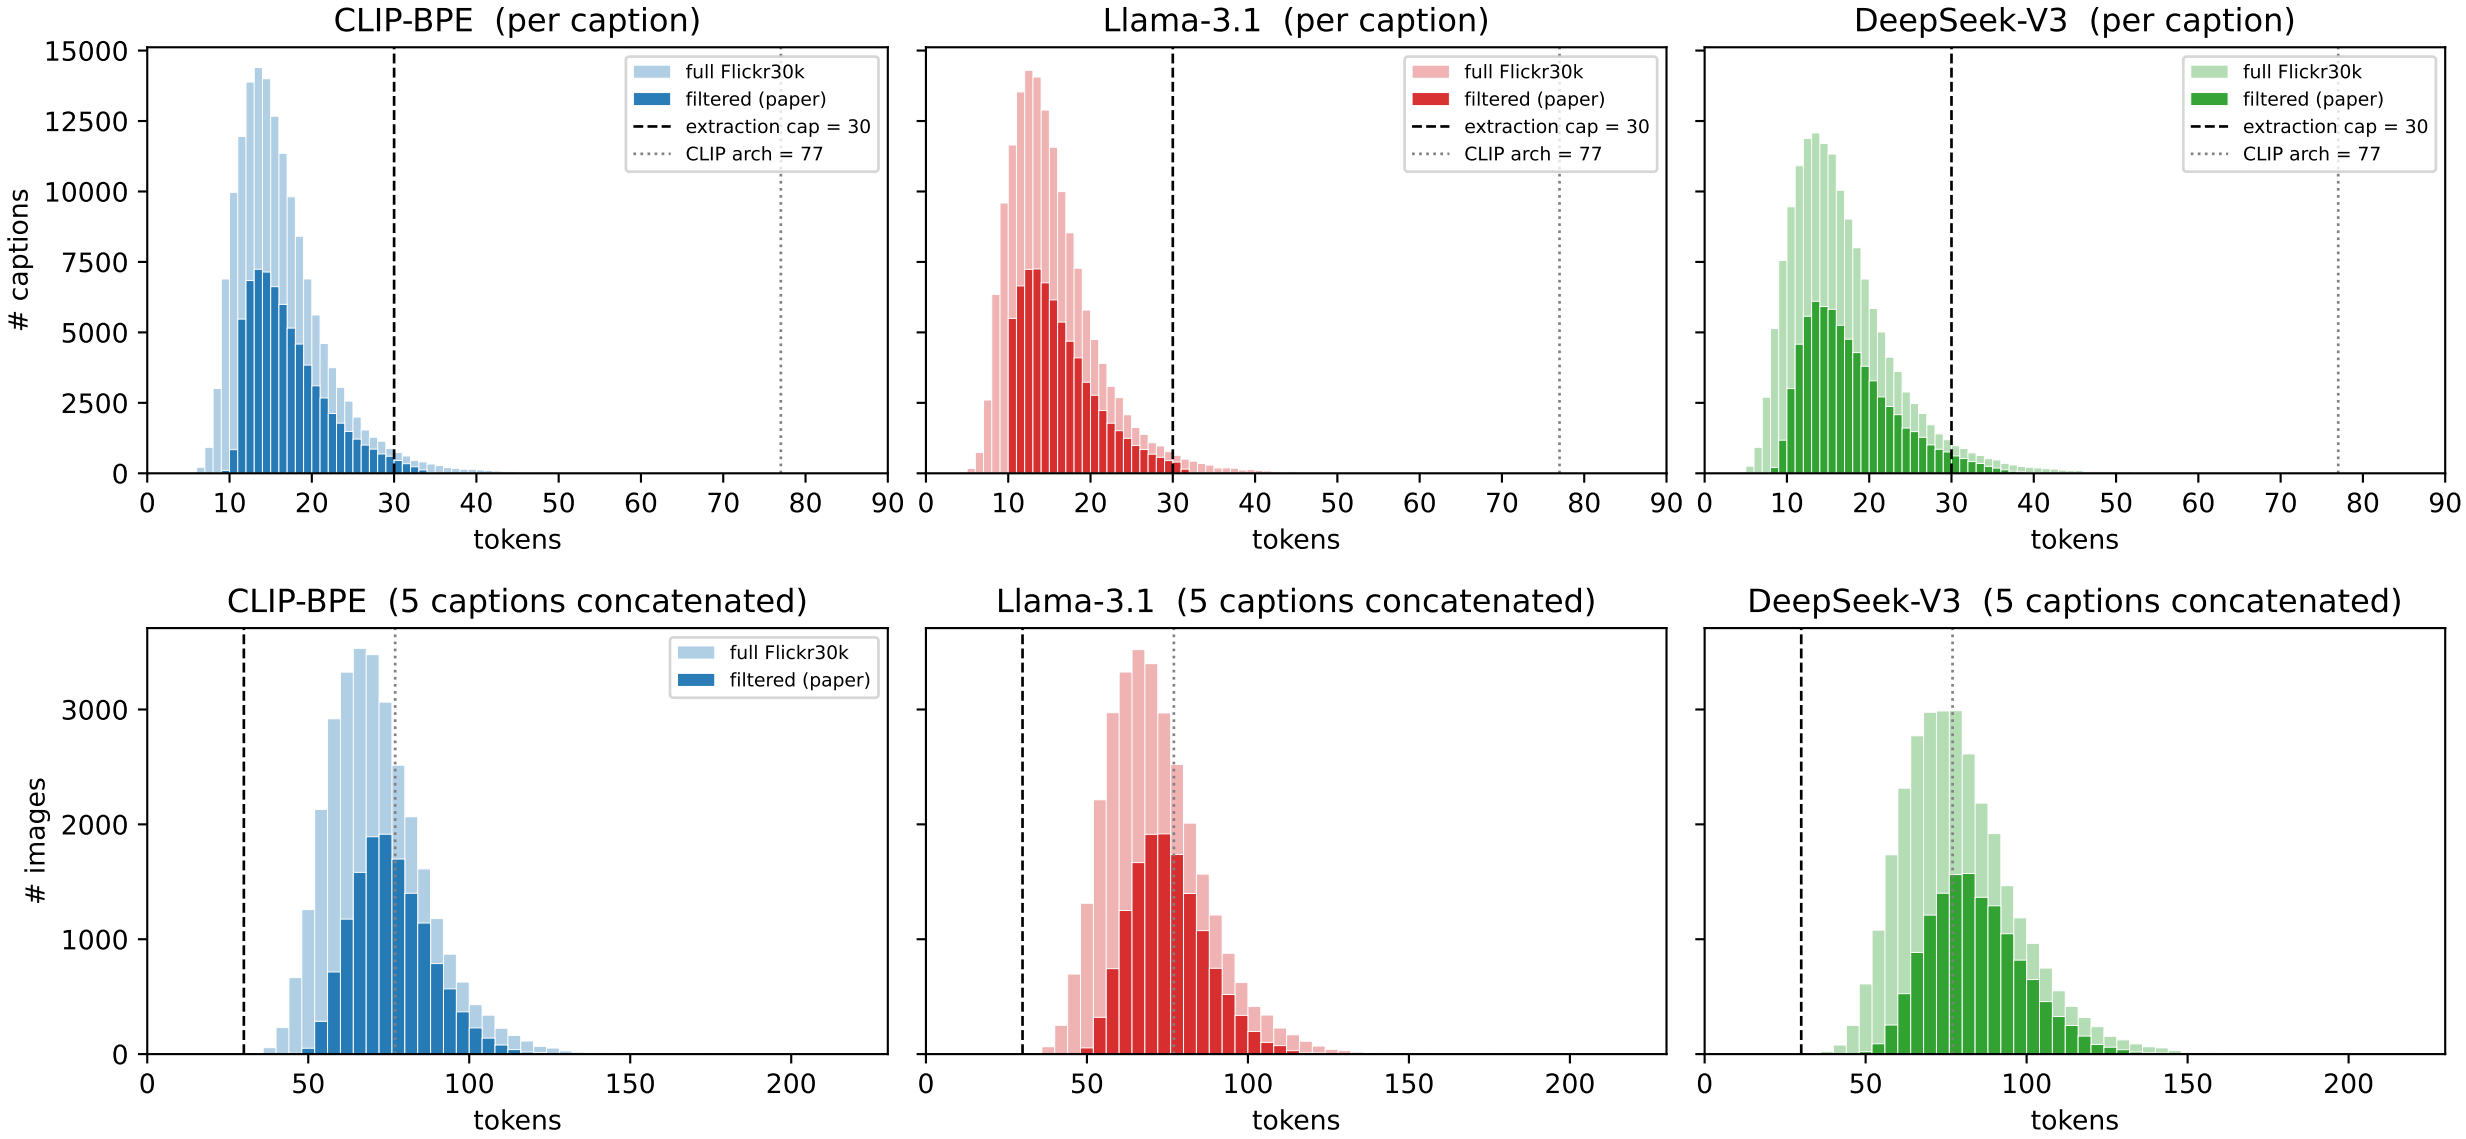
Flickr30k token-length distribution: full vs filtered subset (10–30 Llama-BPE per caption)
CLIP-BPE (per caption)
Llama-3.1 (per caption)
DeepSeek-V3 (per caption)
full Flickr30k
filtered (paper)
extraction cap = 30
CLIP arch = 77
captions
tokens
CLIP-BPE (5 captions concatenated)
Llama-3.1 (5 captions concatenated)
DeepSeek-V3 (5 captions concatenated)
images